

Homework 1

AUTHOR

Adam Whiteside

[Link to the Github repository](#)

Due: Fri, Jan 26, 2024 @ 11:59pm

Please read the instructions carefully before submitting your assignment.

1. This assignment requires you to:

- Upload your Quarto markdown files to a `git` repository
- Upload a `PDF` file on Canvas

2. Don't collapse any code cells before submitting.

3. Remember to make sure all your code output is rendered properly before uploading your submission.

⚠ Please add your name to the the author information in the frontmatter before submitting your assignment.

Question 1

20 points

In this question, we will walk through the process of *forking* a `git` repository and submitting a *pull request*.

1. Navigate to the Github repository [here](#) and fork it by clicking on the icon in the top right



Provide a sensible name for your forked repository when prompted.

2. Clone your Github repository on your local machine

```
$ git clone <<insert your repository url here>>
$ cd hw-1
```

3. In order to activate the `R` environment for the homework, make sure you have `renv` installed beforehand. To activate the `renv` environment for this assignment, open an instance of the `R` console from within the directory and type

```
renv::activate()
```

Follow the instructions in order to make sure that `renv` is configured correctly.

4. Work on the *remaining part* of this assignment as a `.qmd` file.
- Create a `PDF` and `HTML` file for your output by modifying the YAML frontmatter for the Quarto `.qmd` document
5. When you're done working on your assignment, push the changes to your github repository.
6. Navigate to the original Github repository [here](#) and submit a pull request linking to your repository.

Remember to **include your name** in the pull request information!

If you're stuck at any step along the way, you can refer to the [official Github docs here](#)

Question 2

30 points

Consider the following vector

```
my_vec <- c(
  "+0.07",
  "-0.07",
  "+0.25",
  "-0.84",
  "+0.32",
  "-0.24",
  "-0.97",
  "-0.36",
  "+1.76",
```

```
"-0.36"  
)
```

For the following questions, provide your answers in a code cell.

1. What data type does the vector contain?

```
typeof(my_vec)
```

```
[1] "character"
```

2. Create two new vectors called `my_vec_double` and `my_vec_int` which converts `my_vec` to Double & Integer types, respectively,

```
my_vec_double <- as.numeric(my_vec)  
my_vec_int <- as.integer(my_vec)  
my_vec_double
```

```
[1] 0.07 -0.07 0.25 -0.84 0.32 -0.24 -0.97 -0.36 1.76 -0.36
```

```
my_vec_int
```

```
[1] 0 0 0 0 0 0 0 0 1 0
```

3. Create a new vector `my_vec_bool` which comprises of:

- `TRUE` if an element in `my_vec_double` is ≤ 0
- `FALSE` if an element in `my_vec_double` is ≥ 0

How many elements of `my_vec_double` are greater than zero?

```
my_vec_bool <- my_vec_double <= 0  
my_vec_bool
```

```
[1] FALSE TRUE FALSE TRUE FALSE TRUE TRUE TRUE FALSE TRUE
```

```
print(sum(my_vec_double>=0))
```

```
[1] 4
```

4. Sort the values of `my_vec_double` in ascending order.

```
sorted_my_vec_double <- sort(my_vec_double)  
sorted_my_vec_double
```

```
[1] -0.97 -0.84 -0.36 -0.36 -0.24 -0.07 0.07 0.25 0.32 1.76
```

Question 3

50 points

In this question we will get a better understanding of how R handles large data structures in memory.

1. Provide R code to construct the following matrices:

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix} \quad \text{and} \quad \begin{bmatrix} 1 & 2 & 3 & 4 & 5 & \dots & 100 \\ 1 & 4 & 9 & 16 & 25 & \dots & 10000 \end{bmatrix}$$

∴ {cell}

```
mat1 <- matrix(1:9, nrow = 3, byrow = TRUE)
mat1
```

```
      [,1] [,2] [,3]
[1,]     1     2     3
[2,]     4     5     6
[3,]     7     8     9
```

```
mat2 <- matrix(c(1:100, (1:100)^2), nrow = 2, byrow = TRUE)
mat2
```

```
      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13] [,14]
[1,]     1     2     3     4     5     6     7     8     9     10     11     12     13     14
[2,]     1     4     9    16    25    36    49    64    81    100    121    144    169    196
      [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22] [,23] [,24] [,25] [,26]
[1,]    15    16    17    18    19    20    21    22    23    24    25    26
[2,]   225   256   289   324   361   400   441   484   529   576   625   676
      [,27] [,28] [,29] [,30] [,31] [,32] [,33] [,34] [,35] [,36] [,37] [,38]
[1,]    27    28    29    30    31    32    33    34    35    36    37    38
[2,]   729   784   841   900   961  1024  1089  1156  1225  1296  1369  1444
      [,39] [,40] [,41] [,42] [,43] [,44] [,45] [,46] [,47] [,48] [,49] [,50]
[1,]    39    40    41    42    43    44    45    46    47    48    49    50
[2,]  1521  1600  1681  1764  1849  1936  2025  2116  2209  2304  2401  2500
      [,51] [,52] [,53] [,54] [,55] [,56] [,57] [,58] [,59] [,60] [,61] [,62]
[1,]    51    52    53    54    55    56    57    58    59    60    61    62
```

```

[2,] 2601 2704 2809 2916 3025 3136 3249 3364 3481 3600 3721 3844
     [,63] [,64] [,65] [,66] [,67] [,68] [,69] [,70] [,71] [,72] [,73] [,74]
[1,]    63    64    65    66    67    68    69    70    71    72    73    74
[2,] 3969 4096 4225 4356 4489 4624 4761 4900 5041 5184 5329 5476
     [,75] [,76] [,77] [,78] [,79] [,80] [,81] [,82] [,83] [,84] [,85] [,86]
[1,]    75    76    77    78    79    80    81    82    83    84    85    86
[2,] 5625 5776 5929 6084 6241 6400 6561 6724 6889 7056 7225 7396
     [,87] [,88] [,89] [,90] [,91] [,92] [,93] [,94] [,95] [,96] [,97] [,98]
[1,]    87    88    89    90    91    92    93    94    95    96    97    98
[2,] 7569 7744 7921 8100 8281 8464 8649 8836 9025 9216 9409 9604
     [,99] [,100]
[1,]    99    100
[2,] 9801 10000
...

```

Tip

Recall the discussion in class on how `R` fills in matrices

In the next part, we will discover how knowledge of the way in which a matrix is stored in memory can inform better code choices. To this end, the following function takes an input n and creates an $n \times n$ matrix with random entries.

```

generate_matrix <- function(n){
  return(
    matrix(
      rnorm(n^2),
      nrow=n
    )
  )
}

```

For example:

```
generate_matrix(4)
```

```

      [,1]      [,2]      [,3]      [,4]
[1,] -2.4749589  1.1062355  1.0764567 -0.9585682
[2,] -0.2641417  1.0792693 -1.1317516 -0.5551258
[3,] -0.6468692 -0.4357161  0.5886892 -0.3010608
[4,] -0.4105109 -0.9026261 -0.0160500  0.5158264

```

Let `M` be a fixed 50×50 matrix

```

M <- generate_matrix(50)
mean(M)

```

```
[1] 0.01388828
```

2. Write a function `row_wise_scan` which scans the entries of `M` one row after another and outputs the number of elements whose value is ≥ 0 . You can use the following **starter code**

```
row_wise_scan <- function(x){
  n <- nrow(x)
  m <- ncol(x)

  # Insert your code here
  count <- 0
  for(...){
    for(...){
      if(...){
        count <- count + 1
      }
    }
  }

  return(count)
}
```

```
row_wise_scan <- function(x){
  n <- nrow(x)
  m <- ncol(x)

  # Insert your code here
  count <- 0
  for(i in 1:n){
    for(j in 1:m){
      if(x[i,j] >= 0){
        count <- count + 1
      }
    }
  }

  return(count)
}
result <- row_wise_scan(mat1)
print(result)
```

[1] 9

3. Similarly, write a function `col_wise_scan` which does exactly the same thing but scans the entries of `M` one column after another

```
col_wise_scan <- function(x){
  count <- 0

  ... # Insert your code here
```

```
return(count)
```

```
col_wise_scan <- function(x){
  n <- nrow(x)
  m <- ncol(x)

  # Insert your code here
  count <- 0
  for(i in 1:n){
    for(j in 1:m){
      if(x[j,i] >= 0){
        count <- count + 1
      }
    }
  }

  return(count)
}
result <- row_wise_scan(mat1)
print(result)
```

```
[1] 9
```

You can check if your code is doing what it's supposed to using the function [here](#)¹

4. Between `col_wise_scan` and `row_wise_scan`, which function do you expect to take shorter to run? Why? I expect that the `row_wise_scan` would take shorter, because I would guess that going across the columns over and over again would take time.
5. Write a function `time_scan` which takes in a method `f` and a matrix `M` and outputs the amount of time taken to run `f(M)`

```
time_scan <- function(f, M){
  initial_time <- ... # Write your code here
  f(M)
  final_time <- ... # Write your code here

  total_time_taken <- final_time - initial_time
  return(total_time_taken)

  total_time_taken <- final_time - initial_time
  return(total_time_taken)
}
```

```
time_scan <- function(f, M){
  initial_time <- Sys.time() # Write your code here
  f(M)
  final_time <- Sys.time() # Write your code here
```

```
total_time_taken <- final_time - initial_time
return(total_time_taken)}
```

Provide your output to

```
list(
  row_wise_time = time_scan(row_wise_scan, M),
  col_wise_time = time_scan(row_wise_scan, M)
)
```

```
list(
  row_wise_time = time_scan(row_wise_scan, M),
  col_wise_time = time_scan(row_wise_scan, M)
)
```

```
$row_wise_time
Time difference of 0.0001521111 secs
```

```
$col_wise_time
Time difference of 0.0001449585 secs
```

Which took longer to run?

The row wise time function took longer to run.

6. Repeat this experiment now when:

- M is a 100×100 matrix
- M is a 1000×1000 matrix
- M is a 5000×5000 matrix

```
M <- matrix(1:10000, nrow = 100, byrow = TRUE)
list(
  row_wise_time = time_scan(row_wise_scan, M),
  col_wise_time = time_scan(row_wise_scan, M)
)
```

```
$row_wise_time
Time difference of 0.0005772114 secs
```

```
$col_wise_time
Time difference of 0.0006020069 secs
```

```
M <- matrix(1:(1000*1000), nrow = 1000, byrow = TRUE)
list(
  row_wise_time = time_scan(row_wise_scan, M),
  col_wise_time = time_scan(row_wise_scan, M)
)
```



```
$row_wise_time
Time difference of 0.06198382 secs
```

```
$col_wise_time
Time difference of 0.06782603 secs
```

```
M <- matrix(1:(5000*5000), nrow = 5000, byrow = TRUE)
list(
  row_wise_time = time_scan(row_wise_scan, M),
  col_wise_time = time_scan(row_wise_scan, M)
)
```

```
$row_wise_time
Time difference of 1.482058 secs
```

```
$col_wise_time
Time difference of 1.503071 secs
```

What can you conclude?

I can conclude that the col_wise_time function is faster, shown in the three tests above, although a small difference, the col_wise_time had a shorter run time for all three.

Appendix

Print your R session information using the following command

```
sessionInfo()
```

```
R version 4.2.1 (2022-06-23 ucrt)
Platform: x86_64-w64-mingw32/x64 (64-bit)
Running under: Windows 10 x64 (build 19045)
```

```
Matrix products: default
```

```
locale:
[1] LC_COLLATE=English_United States.utf8
[2] LC_CTYPE=English_United States.utf8
[3] LC_MONETARY=English_United States.utf8
[4] LC_NUMERIC=C
[5] LC_TIME=English_United States.utf8
```

```
attached base packages:
[1] stats      graphics  grDevices  utils      datasets  methods   base
```

```
loaded via a namespace (and not attached):
[1] digest_0.6.29    jsonlite_1.8.0    magrittr_2.0.3    evaluate_0.17
```

```
[5] rlang_1.0.6      stringi_1.7.8    cli_3.3.0        rstudioapi_0.14  
[9] rmarkdown_2.17   tools_4.2.1      stringr_1.4.1    htmlwidgets_1.5.4  
[13] xfun_0.33        yaml_2.3.5       fastmap_1.1.0    compiler_4.2.1  
[17] htmltools_0.5.3  knitr_1.40
```

Footnotes

1. If your code is right, the following code should evaluate to be `TRUE`

```
sapply(1:100, function(i) {  
  x <- generate_matrix(100)  
  row_wise_scan(x) == col_wise_scan(x)  
}) %>% sum == 100
```

